

Introduction to Machine Learning

Lecture 4 Linear Models I

Goker Erdogan
26 – 30 November 2018
Pontificia Universidad Javeriana

Types of learning problems

Supervised learning	Unsupervised learning	Reinforcement learning
<p>Learning from supervision (i.e., teacher)</p> <p>Given inputs and outputs, learn a function that maps input → output</p>	<p>Learning with no supervision</p> <p>Given only inputs, extract some pattern from the data.</p>	<p>Learning from rewards and punishments.</p> <p>Inspired by (operant) conditioning in psychology.</p> <p>The agent acts in an environment over some time.</p> <p>It gets reward/punishment from time to time.</p> <p>Needs to learn a <i>policy</i> to maximize reward.</p>
<p>Examples:</p> <ul style="list-style-type: none">- Recognizing faces- Predicting the sales for a product- Learning to rank search results	<p>Examples:</p> <ul style="list-style-type: none">- Clustering (market segmentation)- Dimensionality reduction (visualization)- Generating celebrity faces	<p>Examples:</p> <ul style="list-style-type: none">- Most robotics applications (e.g. learning to move around)- Playing games (e.g., AlphaGo)

Supervised learning

- Two main types of problems
 - **Regression**
 - Outputs are **real-valued** (quantitative)
 - Examples:
 - Predicting the sales of a product
 - Predicting the value of a home
 - **Classification**
 - Outputs are **discrete** (qualitative)
 - Examples:
 - Face/speech/object recognition
 - Predicting if a person will default on their loan
 - Predicting whether a stock will go up or down
- *Note inputs can be of any type (real-valued or discrete).*

Linear regression

- Assume the relation between input and output is **linear**

$$y = \beta_0 + \sum_{d=1}^D x_d \beta_d$$

- β are known as **parameters** or **coefficients**
- β_0 is known as **bias** or **intercept**

- Matrix notation

- $X_{N \times (D+1)}$: Data matrix. N samples. D features (measurements)
 - Add a new feature (column) that is all 1s.
- Y_N : Output (target) vector. N samples.
- β_D : Parameter vector. P coefficients.
- Then, the above can be written as

$$Y = X\beta$$

Linear regression

- Note that we can apply any (possibly **nonlinear**) preprocessing on x

$$y = \beta_0 + \sum_{d=1}^D \phi(x_d) \beta_d$$

- This is still a **linear** model. Because it is **linear in the parameters**.
- Linearity

$$y(c_1\beta + c_2\gamma) = c_1f(\beta) + c_2f(\gamma)$$

- Why do we care about linear regression?
 - Easy to interpret
 - May outperform complex models if
 - Little data
 - High noise in data
 - Foundation for many techniques

Solving linear regression

- Find the parameters that fit the training data $\{x, y\}_{n=1,2,\dots,N}$
 - Pick the *line* that passes as close as possible to the training points
- How do we measure **closeness**?

$$E(\beta) = \frac{1}{2} \sum_{n=1}^N (y - (\beta_0 + \sum_{d=1}^D X_{nd}\beta_d))^2$$

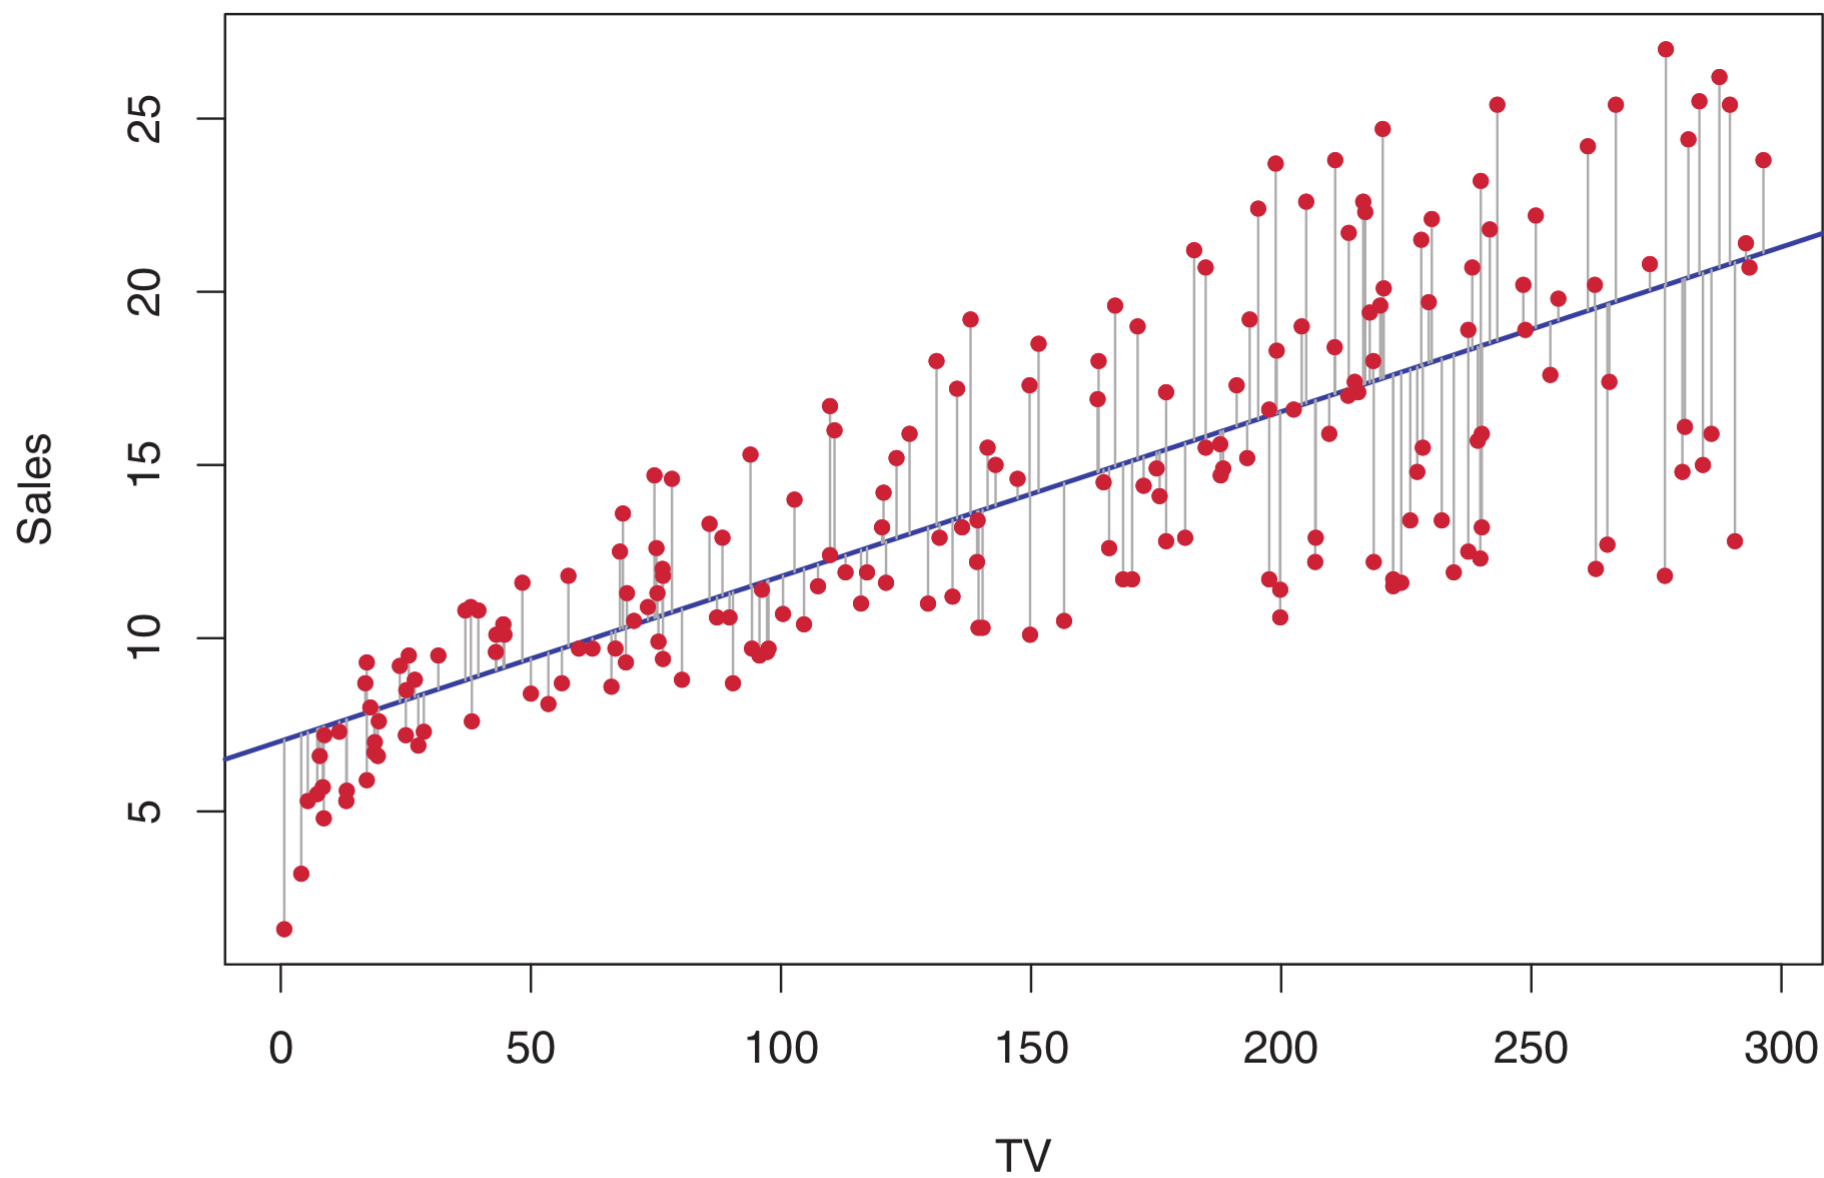
- Find parameters that minimize the **sum-of-squares** error (method of least squares)
 - Set derivative to 0 and solve for β

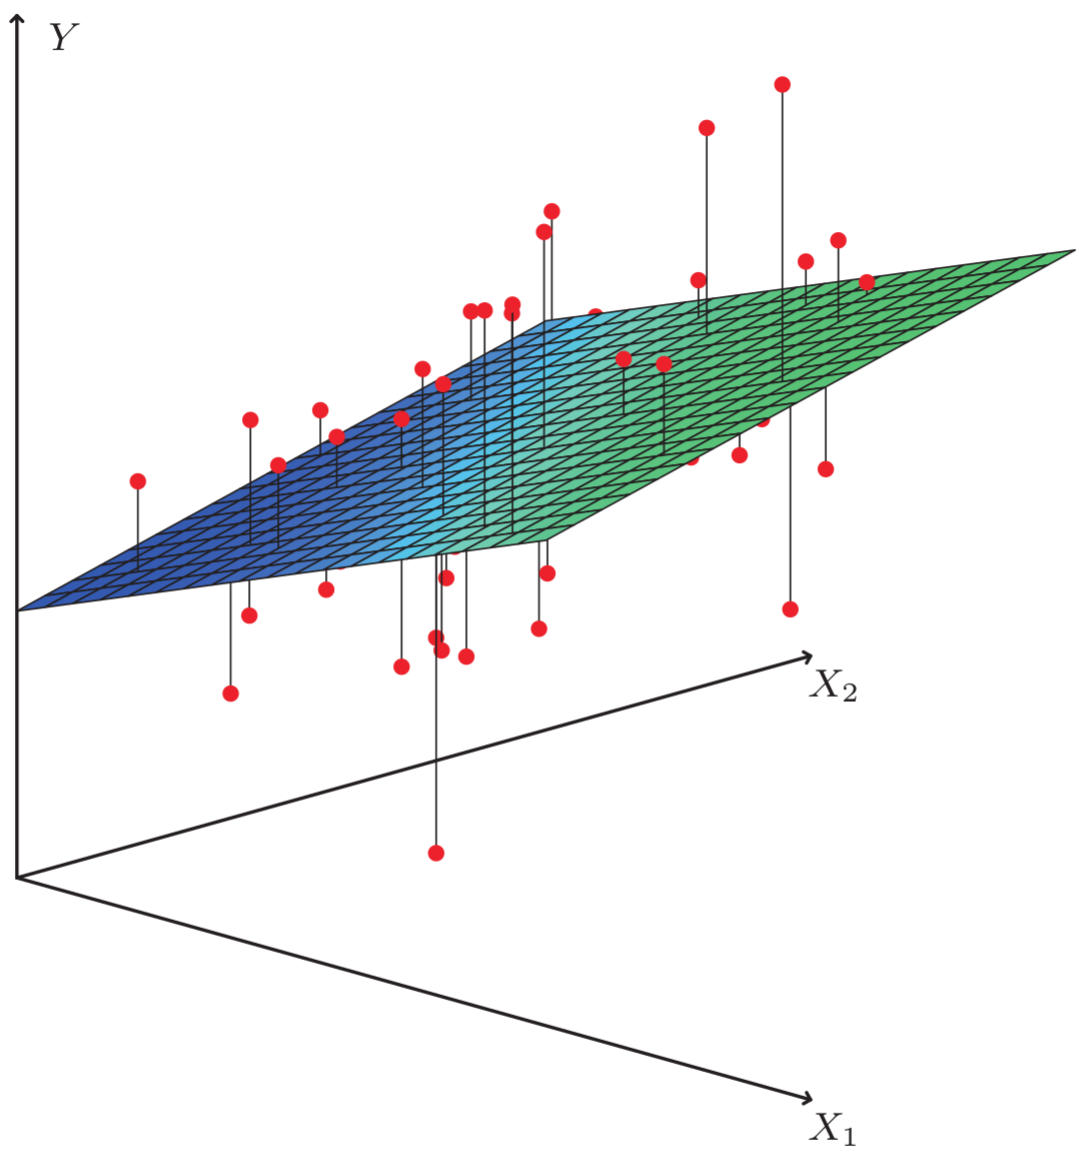
- For 1D case

$$\hat{\beta}_1 = \frac{\sum_{i=1}^n (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^n (x_i - \bar{x})^2},$$

$$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x},$$

\bar{x} : Mean of x
 \bar{y} : Mean of y





Solving linear regression

- Linear regression is an easy problem
 - Convex (i.e., single global minimum)
 - Fast algorithms available

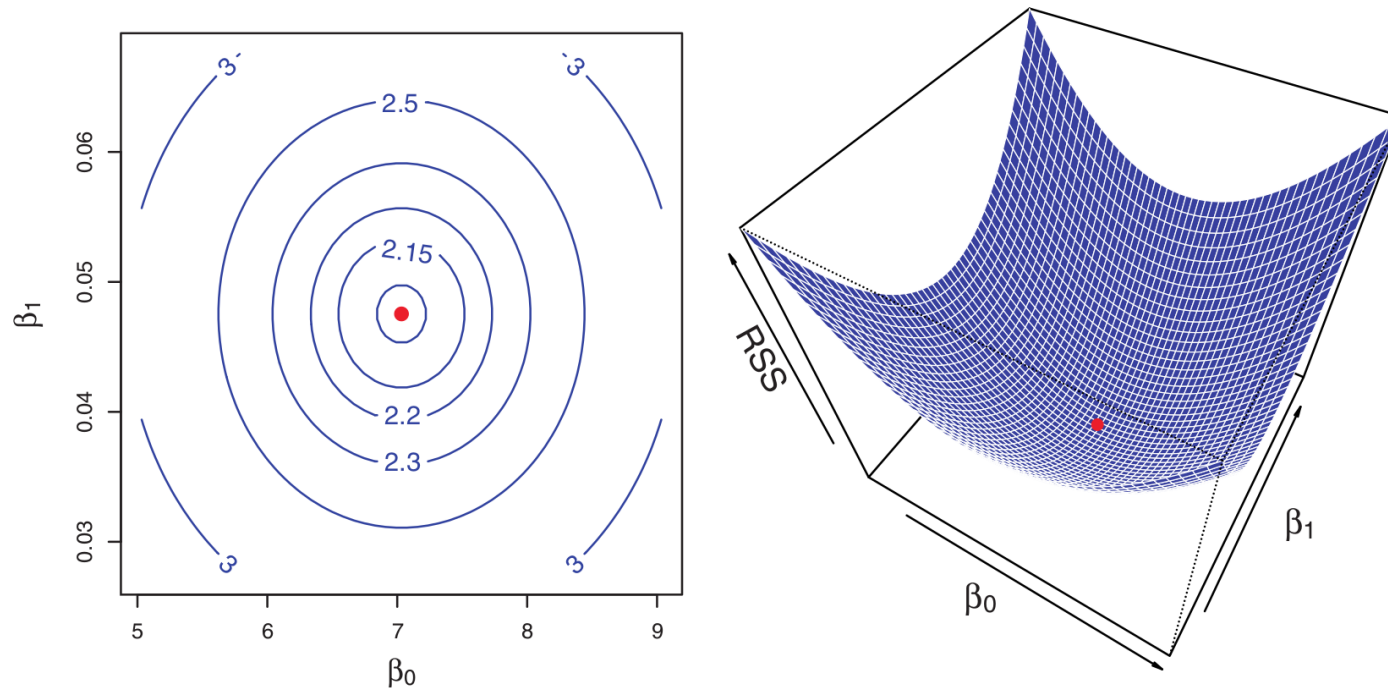


FIGURE 3.2. Contour and three-dimensional plots of the RSS on the **Advertising** data, using **sales** as the response and **TV** as the predictor. The red dots correspond to the least squares estimates $\hat{\beta}_0$ and $\hat{\beta}_1$, given by (3.4).

Solving the general case

- Minimize the following wrt to β

$$\text{RSS}(\beta) = (\mathbf{y} - \mathbf{X}\beta)^T (\mathbf{y} - \mathbf{X}\beta).$$



- Need to assume that \mathbf{X} is full rank (no linearly dependent columns)
- The solution is

$$\hat{\beta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}.$$

- In python, `numpy.linalg.lstsq`
- In R, `fit <- lm(y ~ x1 + x2 + x3 ...)`
- Good resources for matrix calculus
 - <http://www.matrixcalculus.org/>
 - <https://www.math.uwaterloo.ca/~hwolkowi/matrixcookbook.pdf>

What can go wrong?

- Outliers



- Robust alternatives
 - Use absolute value instead of squared error
 - Harder to solve

What can go wrong?

- Correlated inputs (**colinearity**)




FIGURE 3.14. Scatterplots of the observations from the **Credit** data set. Left: A plot of **age** versus **limit**. These two variables are not collinear. Right: A plot of **rating** versus **limit**. There is high collinearity.

What can go wrong?

- Correlated inputs (**colinearity**)
 - Solution unstable
 - Hard to interpret
 - What can we do?
 - Look at the correlation matrix
 - Whiten input data (PCA)
- Correlation of error-terms
 - Violates the independence assumption
 - e.g., Time-series data

What can go wrong?

- Nonlinear input – output relation
 - Interaction terms
 - Residual (error) plot

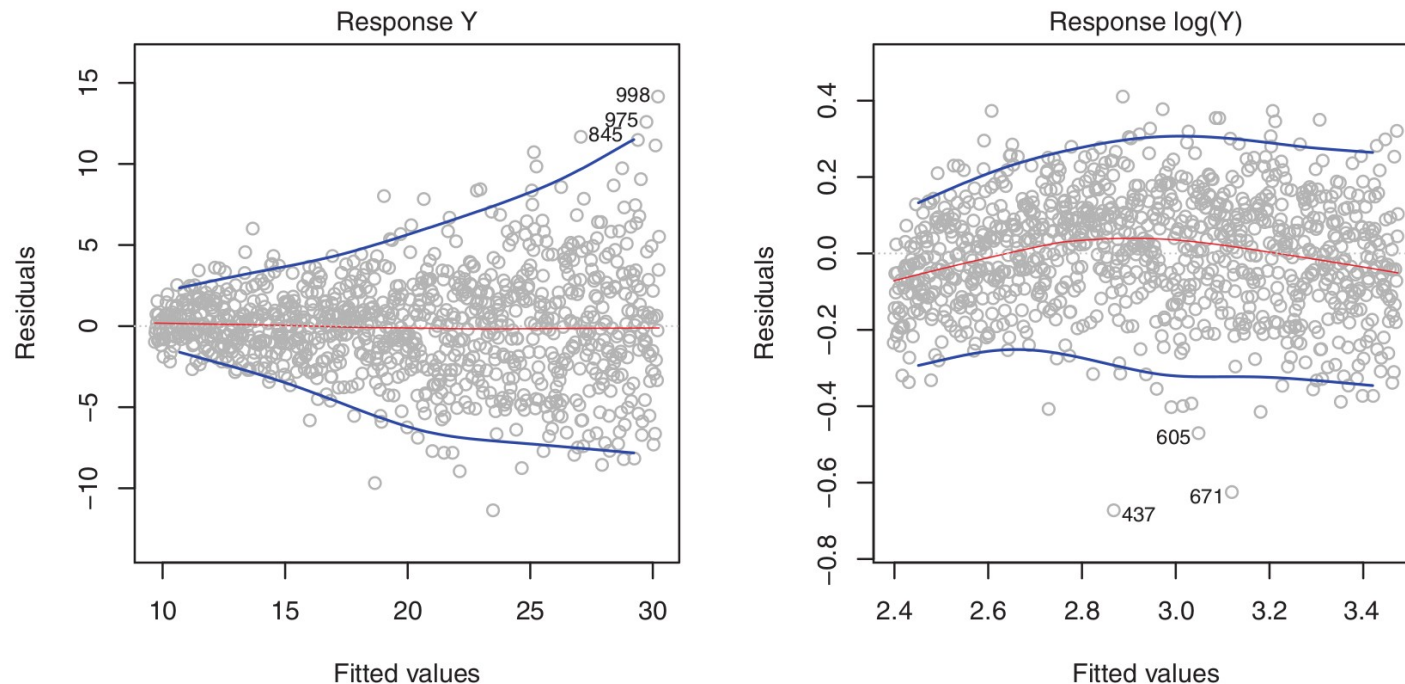


FIGURE 3.11. Residual plots. In each plot, the red line is a smooth fit to the residuals, intended to make it easier to identify a trend. The blue lines track the outer quantiles of the residuals, and emphasize patterns. Left: The funnel shape indicates heteroscedasticity. Right: The response has been log transformed, and there is now no evidence of heteroscedasticity.

Summary

- Linear regression
 - Formulation
 - How to solve it
 - What can go wrong
- Exercises
 - Solve 1D case
 - Do the lab in Section 3.6 of ISLR

References

- [1] James, Witten, Hastie, and Tibshirani. An Introduction to Statistical Learning with Applications in R. Chapter 3.
- [2] Hastie, Tibshirani, and Friedman. The Elements of Statistical Learning. Chapter 3.
- [3] <http://r-statistics.co/Outlier-Treatment-With-R.html>