



Introduction to Machine Learning

Lecture 8 Tree-based models I

Goker Erdogan
26 – 30 November 2018
Pontificia Universidad Javeriana


Decision trees

- A **popular technique** for supervised learning
 - Both for regression and classification
- **Idea**
 - Split the input space into rectangular regions
 - Predict using a constant value for each region



Decision trees

- Advantages
 - Easy to interpret
 - Easy to display graphically
 - Non-linear
- Terminology
 - Internal nodes
 - Branches
 - Leaves or **terminal nodes**



Regression trees

- Given N training samples of $\{x, y\}_{n=1\dots N}$
- Split the input space into rectangles
- For each leaf (region)
 - prediction = the mean of the target values in that region
 - Why? Because that is the prediction that minimizes error (sum-of-squares).






FIGURE 8.1. For the **Hitters** data, a regression tree for predicting the log salary of a baseball player, based on the number of years that he has played in the major leagues and the number of hits that he made in the previous year. At a given internal node, the label (of the form $X_j < t_k$) indicates the left-hand branch emanating from that split, and the right-hand branch corresponds to $X_j \geq t_k$. For instance, the split at the top of the tree results in two large branches. The left-hand branch corresponds to **Years**<4.5, and the right-hand branch corresponds to **Years**>=4.5. The tree has two internal nodes and three terminal nodes, or leaves. The number in each leaf is the mean of the response for the observations that fall there.

Fitting a regression tree

- Regression tree with M leaves (regions)

- Given x , our prediction is $f(x)$

$$f(x) = \sum_{m=1}^M \bar{y}_m I(x \in R_m)$$

- where \bar{y}_m is the **average** of all y in R_m


- Given a training set $\{x, y\}_{n=1\dots N}$

- Find **the partitioning** of space (i.e., rectangles) that **minimizes error**


$$RSS = \sum_m \sum_{i \in R_m} (y_i - \bar{y}_m)^2$$

Fitting a regression tree

- Why not consider **all possible partitionings**?
 - The number of possible partitionings is huge!



- We consider only (recursive) **binary splits on single variables**



Fitting a regression tree

- Various algorithms
 - CART
 - C4.5 (C5.0)
- These are usually
 - **Top-down**: Start from all data at the root
 - **Greedy**: At each point, pick the split that maximizes reduction in error

General outline for regression tree fitting procedure

- For each **feature j**
 - Find the best **split position t**
 - Calculate the reduction in RSS
- Pick the **feature j** that reduces the error the most
- Add a new node splitting **feature j at value t**
- Repeat until reached a stopping criteria

Fitting a regression tree

- How do we find the **best split point t** for a feature?

$$R_1(j, t) = \{X | X_j < t\} \text{ and } R_2(j, t) = \{X | X_j \geq t\}$$

$$\min_t \sum_{i: x_i \in R_1} (y_i - \bar{y}_{R_1})^2 + \sum_{i: x_i \in R_2} (y_i - \bar{y}_{R_2})^2$$

- No easy method
- **Try every possible** (or a subset of all) **t**
 - Look at the training data, **find all unique y values** and sort them
 - **Try each y** one by one and calculate the error
 - Pick the one that has **the minimum error**

When do we stop?

- Popular **stopping criteria**
 - Stop when there are **fewer than K samples** in each region
 - Stop when tree reaches a **maximum depth**
- Another idea
 - Split if only reduction in $RSS > R$
 - Empirically not a good rule
 - Too **short-sighted**
- We don't want a very large tree
 - Risk of **overfitting**
 - Think of **tree size as a complexity parameter**
- Build a large tree then **prune it**

How to prune the tree

- Weakest link pruning
 - For each internal node
 - Calculate the increase in error if we removed that node
 - Remove the node with smallest increase in error
 - Evaluate the new tree on a validation set
 - Repeat
- Note we evaluate on a validation set
 - e.g., use k-fold cross-validation




FIGURE 8.5. Regression tree analysis for the **Hitters** data. The training, cross-validation, and test MSE are shown as a function of the number of terminal nodes in the pruned tree. Standard error bands are displayed. The minimum cross-validation error occurs at a tree size of three.



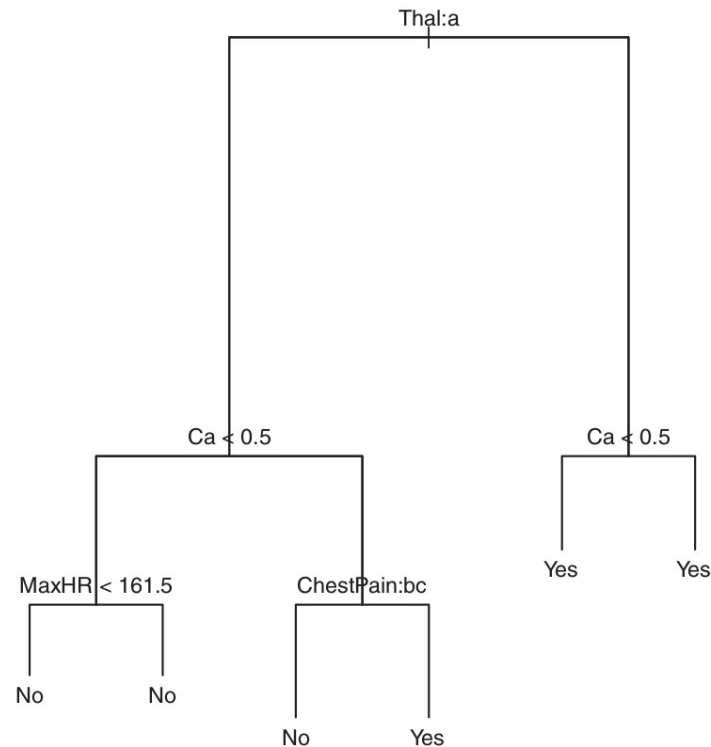
FIGURE 8.1. For the **Hitters** data, a regression tree for predicting the log salary of a baseball player, based on the number of years that he has played in the major leagues and the number of hits that he made in the previous year. At a given internal node, the label (of the form $X_j < t_k$) indicates the left-hand branch emanating from that split, and the right-hand branch corresponds to $X_j \geq t_k$. For instance, the split at the top of the tree results in two large branches. The left-hand branch corresponds to **Years**<4.5, and the right-hand branch corresponds to **Years**>=4.5. The tree has two internal nodes and three terminal nodes, or leaves. The number in each leaf is the mean of the response for the observations that fall there.

Classification tree

- **Classification tree** with M leaves (regions)
 - Given x, our prediction is f(x)

$$f(x) = \sum_{m=1}^M k_m I(x \in R_m)$$

- where k_m is the **majority** class in R_m



How to split nodes

- We want to **minimize misclassification rate**

$$R_1(j, t) = \{X | X_j < t\} \text{ and } R_2(j, t) = \{X | X_j \geq t\}$$

$$\min_t N_{R_1} E(R_1, k) + N_{R_2} E(R_2, k)$$

- N_{R_i} is the **number of samples** falling in region R_i
- $E(R, k)$ is some **error function**
- What can E be?

- **Misclassification rate:** $1 - \max_k (\bar{p}_{R,k})$

$\bar{p}_{R,k} =$ Ratio of points
from class k in
region R

- **Gini index:** $\sum_k \bar{p}_{R,k} (1 - \bar{p}_{R,k})$

- **Entropy:** $-\sum_k \bar{p}_{R,k} \log(\bar{p}_{R,k})$


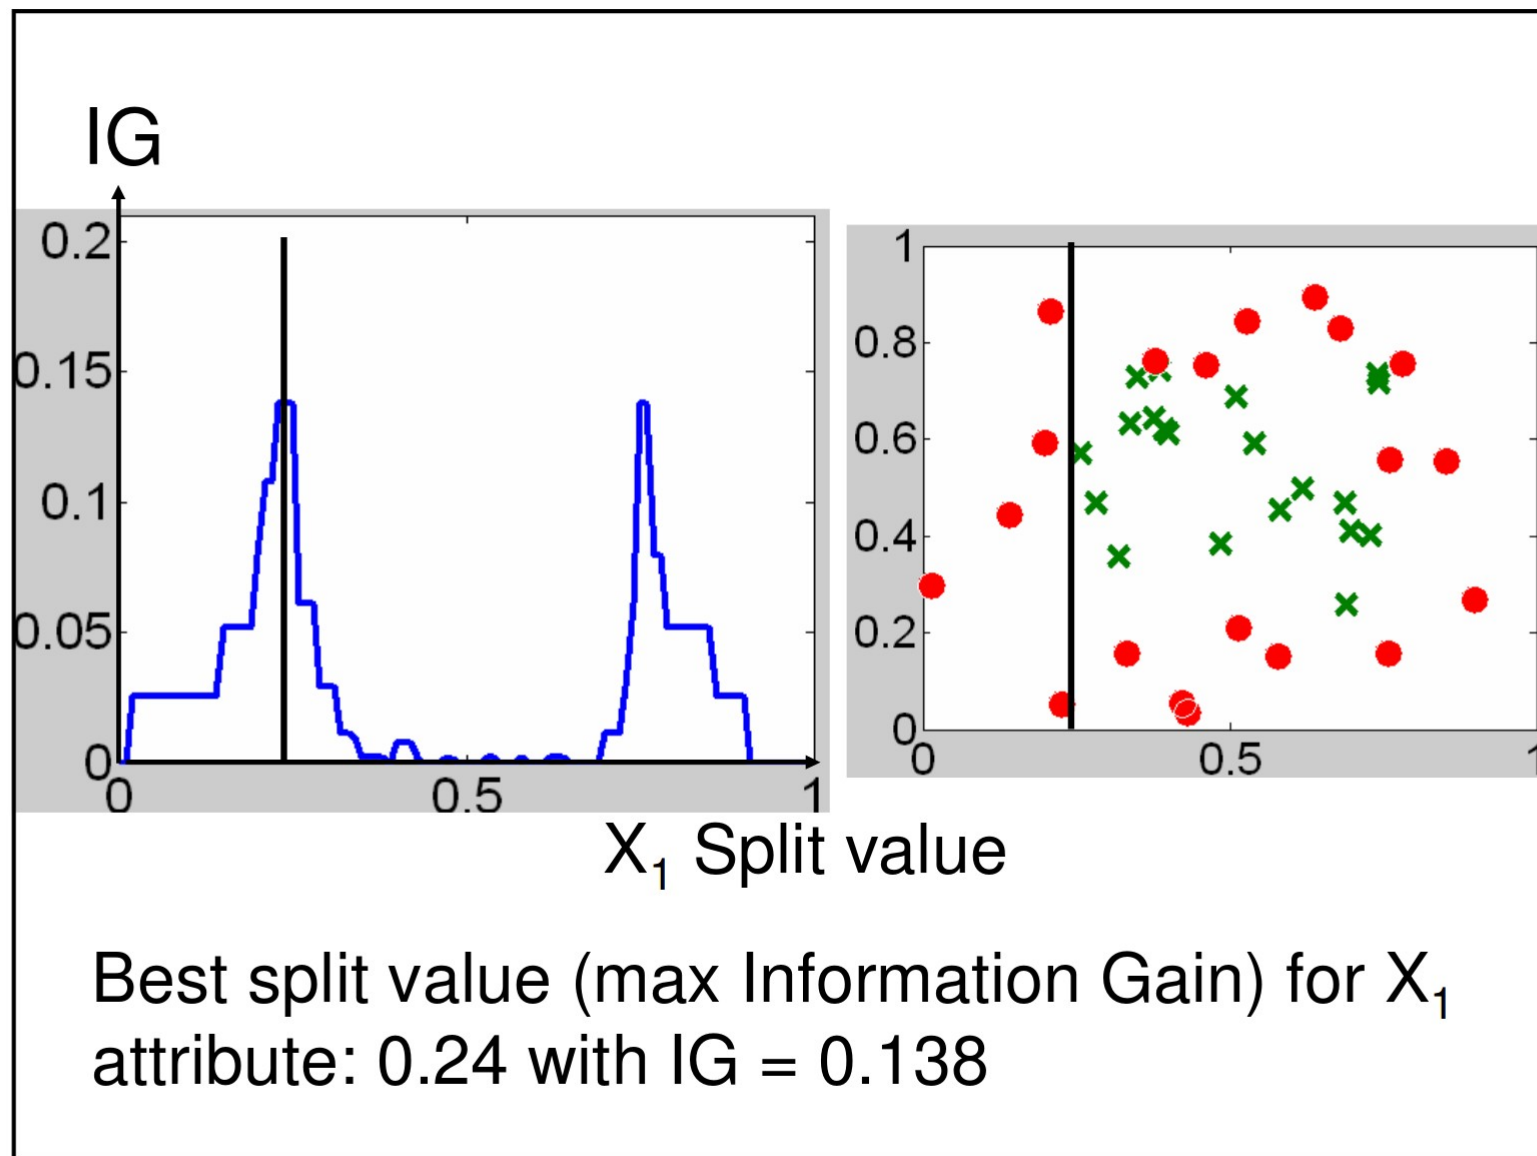
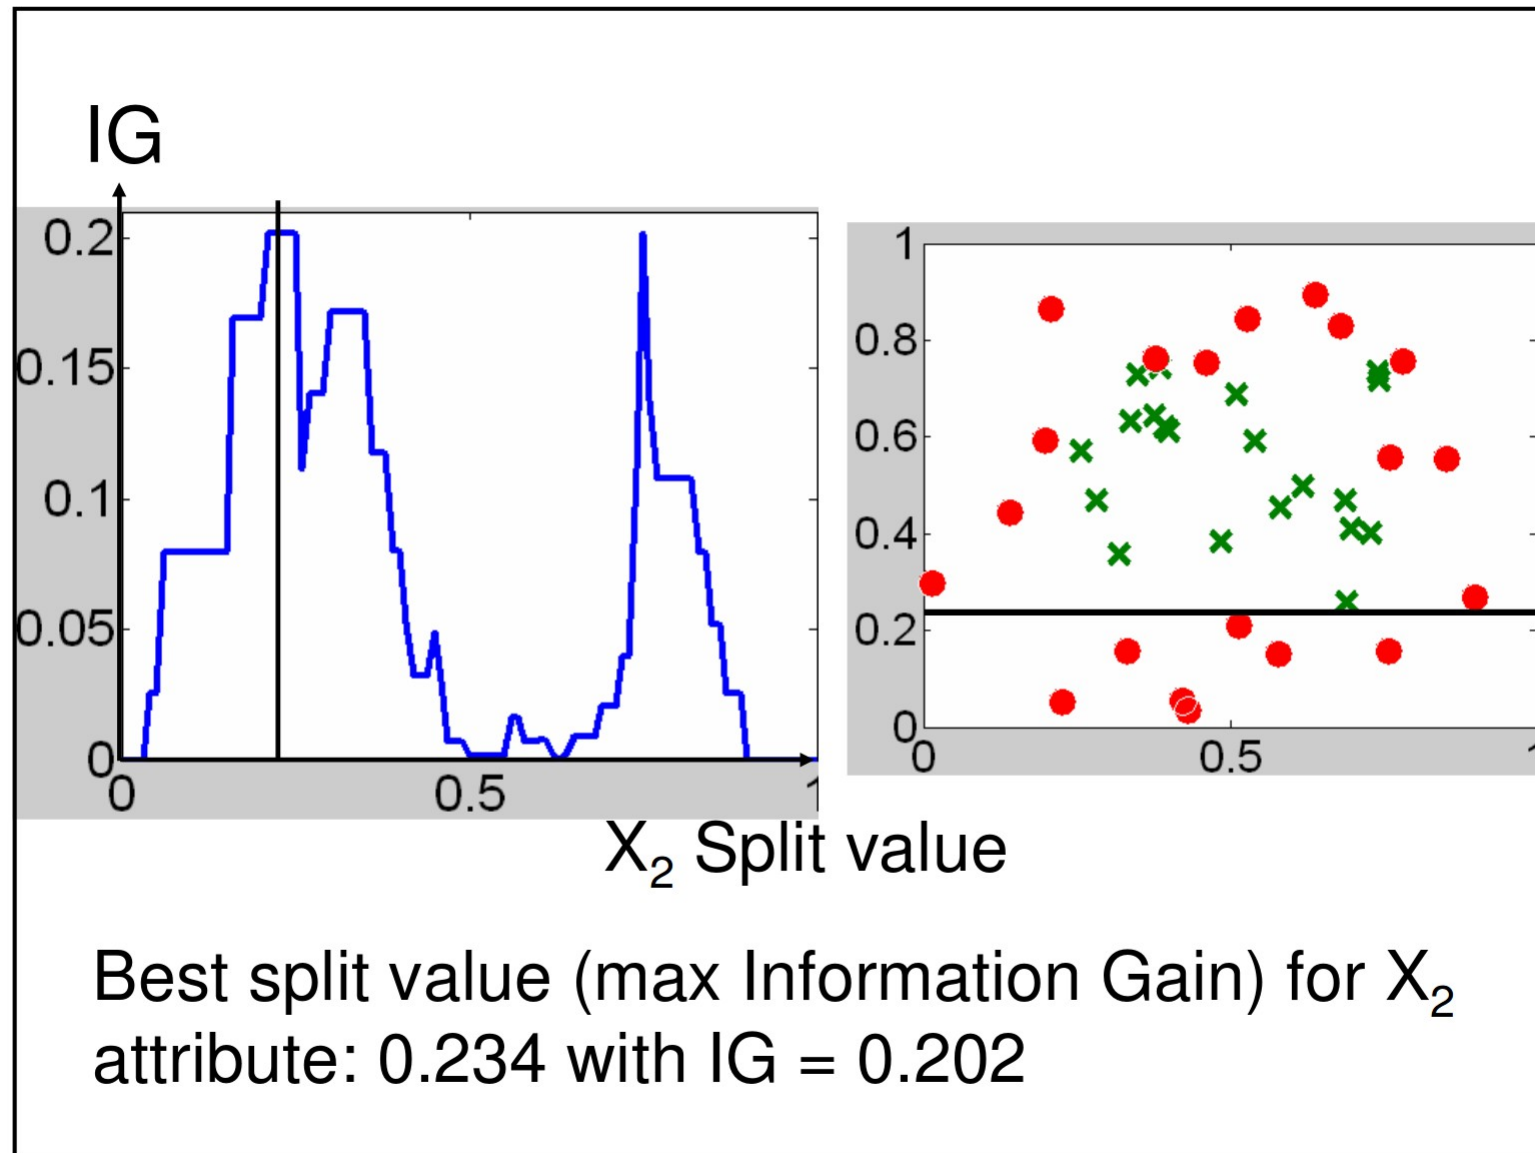


FIGURE 9.3. Node impurity measures for two-class classification, as a function of the proportion p in class 2. Cross-entropy has been scaled to pass through $(0.5, 0.5)$.

Why Gini or entropy?

- Gini and entropy **works better** in practice
- They care more about **purity**
 - Split 1: Left \rightarrow (300, 100), Right \rightarrow (100, 300)
 - Split 2: Left \rightarrow (200, 400), Right \rightarrow (200, 0)
 - Misclassification rate: 0.25 for both
 - But **split 2 is purer**
- Gini and entropy are **differentiable**
- Classification tree fitting
 - Use the **same algorithm** as regression trees
 - Using Gini/entropy instead of RSS to split nodes
 - Prune
 - Use misclassification rate to do the pruning





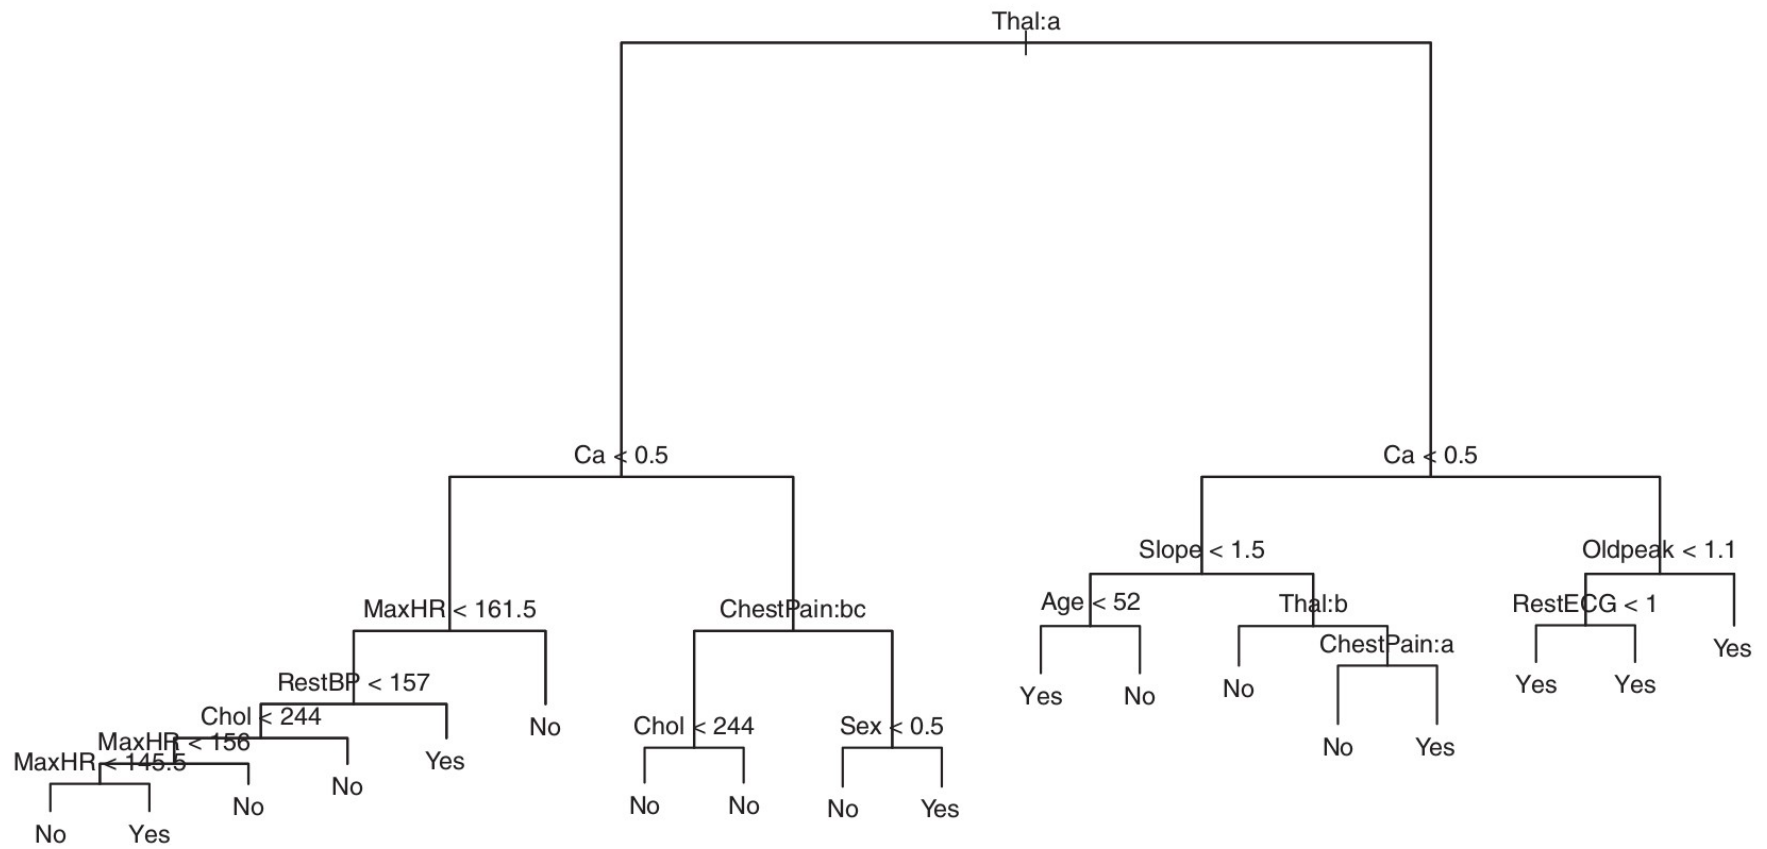


FIGURE 8.6. Heart data. Top: The unpruned tree.





FIGURE 8.6. *Heart* data. Top: The unpruned tree. Bottom Left: Cross-validation error, training, and test error, for different sizes of the pruned tree. Bottom Right: The pruned tree corresponding to the minimal cross-validation error.

Advantages/disadvantages

- Advantages
 - Easy to **interpret**
 - Easy to display graphically
 - **Non-linear**
 - Can handle **qualitative predictors**
 - Can handle **missing data**
- Disadvantages
 - Does **not perform very well**
 - High variance, mostly because the tree is built sequentially
 - Various improvements: bagging/boosting/random forests



Disadvantages

- Non-smooth



Disadvantages

- Assumption of rectangular regions might be problematic



Summary

- Decision trees
- Regression trees
 - Fitting algorithm
 - Stopping criteria
 - Pruning
- Classification trees
 - Splitting criteria: Gini, entropy
- Advantages/disadvantages
- Exercises
 - Do the labs in Section 8.3.1 and 8.3.2 in ISLR

References

- [1] James, Witten, Hastie, and Tibshirani. An Introduction to Statistical Learning with Applications in R. Chapter 8.
- [2] Hastie, Tibshirani, and Friedman. The Elements of Statistical Learning. Chapter 9.
- [3] Bishop. Pattern Recognition and Machine Learning. Chapter 14.
- [4] Murphy. Machine Learning: A Probabilistic Perspective. Chapter 16.
- [5] <https://www.cs.cmu.edu/afs/cs.cmu.edu/academic/class/15381-s06/www/DTs.pdf>